


A guided tour in targeted learning territory

David Benkeser, Antoine Chambaz, Nima Hejazi

08/06/2018

Contents

1	Introduction	1
2	A simulation study	1
2.1	Reproducible experiment as a law.	1
2.2	The parameter of interest, first pass.	3
2.3	The parameter of interest, second pass.	4
2.4	Being smooth, first pass.	6
2.5	 Being smooth, second pass.	7
2.6	The efficient influence curve.	7
2.7	Computing and comparing Cramér-Rao bounds.	8
2.8	Revisiting Section 2.4.	8
2.9	Double-robustness	9
2.10	Targeted inference.	10

1 Introduction

This is a very first draft of our article. The current **tentative** title is "A guided tour in targeted learning territory".

Explain our objectives and how we will meet them. Explain that the symbol  indicates more delicate material.

Use sectioning a lot to ease cross-referencing.

Do we include exercises?

```
set.seed(54321) ## because reproducibility matters...
suppressMessages(library(R.utils)) ## make sure it is installed
suppressMessages(library(tidyverse)) ## make sure it is installed
suppressMessages(library(ggplot2)) ## make sure it is installed
expit <- plogis
logit <- qlogis
```

2 A simulation study

blabla

2.1 Reproducible experiment as a law. We are interested in a reproducible experiment. The generic summary of how one realization of the experiment unfolds, our observation, is called O . We view O as a random variable drawn from what we call the law P_0 of the experiment. The law P_0 is viewed as an element of what we call the model. Denoted by \mathcal{M} , the model is the collection of *all* laws from which O can be drawn

and that meet some constraints. The constraints translate the knowledge we have about the experiment. The more we know about the experiment, the smaller is \mathcal{M} . In all our examples, model \mathcal{M} will put very few restrictions on the candidate laws.

Consider the following chunk of code:

```
draw_from_experiment <- function(n, full = FALSE) {
  ## preliminary
  n <- Arguments$getInteger(n, c(1, Inf))
  full <- Arguments$getLogical(full)
  ## ## 'Gbar' and 'Qbar' factors
  Gbar <- function(W) {
    expit(-0.2 + 3 * sqrt(W) - 1.5 * W)
  }
  Qbar <- function(AW) {
    A <- AW[, 1]
    W <- AW[, 2]
    A * cos((1 + W) * pi / 5) + (1 - A) * sin((1 + W^2) * pi / 4)
  }
  ## sampling
  ## ## context
  W <- runif(n)
  ## ## counterfactual rewards
  zeroW <- cbind(A = 0, W)
  oneW <- cbind(A = 1, W)
  Qbar.zeroW <- Qbar(zeroW)
  Qbar.oneW <- Qbar(oneW)
  Yzero <- rbeta(n, shape1 = 1, shape2 = (1 - Qbar.zeroW) / Qbar.zeroW)
  Yone <- rbeta(n, shape1 = 1, shape2 = (1 - Qbar.oneW) / Qbar.oneW)
  ## ## action undertaken
  A <- rbinom(n, size = 1, prob = Gbar(W))
  ## ## actual reward
  Y <- A * Yone + (1 - A) * Yzero
  ## ## observation
  if (full) {
    obs <- cbind(W = W, Yzero = Yzero, Yone = Yone, A = A, Y = Y)
  } else {
    obs <- cbind(W = W, A = A, Y = Y)
  }
  attr(obs, "Gbar") <- Gbar
  attr(obs, "Qbar") <- Qbar
  attr(obs, "QW") <- dunif
  ##
  return(obs)
}
```

We can interpret `draw_from_experiment` as a law P_0 since we can use the function to sample observations from a common law. It is even a little more than that, because we can tweak the experiment, by setting its `full` argument to `TRUE`, in order to get what appear as intermediary (counterfactual) variables in the regular experiment. The next chunk of code runs the (regular) experiment five times independently:

```
(five_obs <- draw_from_experiment(5))
```

```
##           W A           Y
## [1,] 0.4290078 0 0.9998700
## [2,] 0.4984304 1 0.9351501
```

```
## [3,] 0.1766923 0 0.9477263
## [4,] 0.2743935 1 0.8287541
## [5,] 0.2165102 1 0.9977092
## attr("Gbar")
## function (W)
## {
##     expit(-0.2 + 3 * sqrt(W) - 1.5 * W)
## }
## <bytecode: 0x4b1ac00>
## <environment: 0x5770f20>
## attr("Qbar")
## function (AW)
## {
##     A <- AW[, 1]
##     W <- AW[, 2]
##     A * cos((1 + W) * pi/5) + (1 - A) * sin((1 + W^2) * pi/4)
## }
## <bytecode: 0x6410dc8>
## <environment: 0x5770f20>
## attr("QW")
## function (x, min = 0, max = 1, log = FALSE)
## .Call(C_dunif, x, min, max, log)
## <bytecode: 0x6403cf8>
## <environment: namespace:stats>
```

We can view the attributes of object `five_obs` because, in this section, we act as oracles, *i.e.*, we know completely the nature of the experiment. From a probabilistic point of view, the attributes `Gbar`, `Qbar` and `QW` are infinite-dimensional features of P_0 . There is more to P_0 than \bar{G}_0 (`Gbar`), \bar{Q}_0 (`Qbar`), formally defined by

$$\bar{G}_0(W) \equiv P_0(A = 1|W), \quad \bar{Q}_0(A, W) \equiv E_{P_0}(Y|A, W), \quad (1)$$

and the marginal distribution $Q_{0,W}$ of W under P_0 (`QW`), for instance the conditional distribution (not expectation) of Y given (A, W) , but \bar{G}_0 , \bar{Q}_0 and $Q_{0,W}$ will play a prominent role in our story.

2.2 The parameter of interest, first pass. It happens that we especially care for a finite-dimensional feature of P_0 that we denote by ψ_0 . Its definition involves the aforementioned infinite-dimensional features:

$$\begin{aligned} \psi_0 &\equiv E_{P_0}(\bar{Q}_0(1, W) - \bar{Q}_0(0, W)) \\ &= \int (\bar{Q}_0(1, w) - \bar{Q}_0(0, w)) dQ_{0,W}(w). \end{aligned} \quad (2)$$

Acting as oracles, we can compute explicitly the numerical value of ψ_0 .

Our interest in ψ_0 is of causal nature. Taking a closer look at `drawFromExperiment` reveals indeed that the random making of an observation O drawn from P_0 can be summarized by the following causal graph and nonparametric system of structural equations:

```
## plot the causal diagram
```

and, for some deterministic functions f_w, f_a, f_y and independent sources of randomness U_w, U_a, U_y ,

1. sample the context where the rest of the experiment will take place, $W = f_w(U_w)$;

2. sample the two counterfactual rewards of the two actions that can be undertaken, $Y_0 = f_y(0, W, U_y)$ and $Y_1 = f_y(1, W, U_y)$;
3. sample which action is carried out in the given context, $A = f_a(W, U_a)$;
4. define the corresponding reward, $Y = AY_1 + (1 - A)Y_0$;
5. summarize the course of the experiment with the observation $O = (W, A, Y)$, thus concealing Y_0 and Y_1 .

The above description of the experiment `draw_from_experiment` is useful to ram home what it means to run the “full” experiment by setting argument `full` to `TRUE` in a call to `draw_from_experiment`. Doing so triggers a modification of the nature of the experiment, enforcing that the counterfactual rewards Y_0 and Y_1 be part of the summary of the experiment eventually. In light of the above enumeration, $\mathbb{O} \equiv (W, Y_0, Y_1, A, Y)$ is output, as opposed to its summary measure O . This defines another experiment and its law, that we denote \mathbb{P}_0 .

It is well known (do we give the proof or refer to other articles?) that

$$\psi_0 = E_{\mathbb{P}_0}(Y_1 - Y_0).$$

Thus, ψ_0 compares (additively) the averages of the two counterfactual rewards. In other words, ψ_0 quantifies the difference in average of the reward one would get in a world where one would always enforce action $a = 1$ with the reward one would get in a world where one would always enforce action $a = 0$. This said, it is worth emphasizing that ψ_0 is a well defined parameter beyond its causal interpretation.

To conclude this subsection, we draw advantage from the possibility to sample full observations from `draw_from_experiment` by setting its argument `full` to `TRUE` in order to numerically approximate ψ_0 . By the law of large numbers, the following chunk of code approximates ψ_0 :

```
B <- 1e6
full_obs <- draw_from_experiment(B, full = TRUE)
(psi_hat <- mean(full_obs[, "Yone"] - full_obs[, "Yzero"]))

## [1] -0.2644049
```

In fact, the central limit theorem and Slutsky’s lemma allow us to build a confidence interval with asymptotic level 95% for ψ_0 :

```
sd_hat <- sd(full_obs[, "Yone"] - full_obs[, "Yzero"])
alpha <- 0.05
(psi_CI <- psi_hat + c(-1, 1) * qnorm(1 - alpha / 2) * sd_hat / sqrt(B))

## [1] -0.2652679 -0.2635419
```

2.3 The parameter of interest, second pass. Suppose we know beforehand that O drawn from P_0 takes its values in $\mathcal{O} \equiv [0, 1] \times \{0, 1\} \times [0, 1]$ and that $P_0(A = 1|W)$ is bounded away from zero and one $Q_{0,W}$ -almost surely (this is the case indeed). Then we can define model \mathcal{M} as the set of all laws P on \mathcal{O} such that $\tilde{G}(W) \equiv P(A = 1|W)$ is bounded away from zero and one Q_W -almost surely, where Q_W is the marginal distribution of W under P .

Let us also define generically \bar{Q} as

$$\bar{Q}(A, W) \equiv E_P(Y|A, W).$$

Central to our approach is viewing ψ_0 as the value at P_0 of the statistical mapping Ψ from \mathcal{M} to $[0, 1]$ characterized by

$$\begin{aligned}\Psi(P) &\equiv E_P(\bar{Q}(1, W) - \bar{Q}(0, W)) \\ &= \int (\bar{Q}(1, w) - \bar{Q}(0, w)) dQ_W(w),\end{aligned}$$

a clear extension of (2). For instance, although the law $\Pi_0 \in \mathcal{M}$ encoded by default (*i.e.*, with $h=0$) in `drawFromAnotherExperiment` defined below differs starkly from P_0 ,

```
draw_from_another_experiment <- function(n, h = 0) {
  ## preliminary
  n <- Arguments$getInteger(n, c(1, Inf))
  h <- Arguments$getNumeric(h)
  ## ## 'Gbar' and 'Qbar' factors
  Gbar <- function(W) {
    sin((1 + W) * pi / 6)
  }
  Qbar <- function(AW, hh = h) {
    A <- AW[, 1]
    W <- AW[, 2]
    expit( logit( A * W + (1 - A) * W^2 ) +
           hh * 10 * sqrt(W) * A )
  }
  ## sampling
  ## ## context
  W <- runif(n, min = 1/10, max = 9/10)
  ## ## action undertaken
  A <- rbinom(n, size = 1, prob = Gbar(W))
  ## ## reward
  shape1 <- 4
  QAW <- Qbar(cbind(A, W))
  Y <- rbeta(n, shape1 = shape1, shape2 = shape1 * (1 - QAW) / QAW)
  ## ## observation
  obs <- cbind(W = W, A = A, Y = Y)
  attr(obs, "Gbar") <- Gbar
  attr(obs, "Qbar") <- Qbar
  attr(obs, "QW") <- function(x){dunif(x, min = 1/10, max = 9/10)}
  attr(obs, "shape1") <- shape1
  ##
  return(obs)
}
```

parameter $\Psi(\Pi_0)$ is well defined, and numerically approximated by `psi.Pi.zero` in the following chunk of code:

```
five_obs_from_another_experiment <- draw_from_another_experiment(5)
integrand <- function(w) {
  Qbar <- attr(five_obs_from_another_experiment, "Qbar")
  QW <- attr(five_obs_from_another_experiment, "QW")
  ( Qbar(cbind(1, w)) - Qbar(cbind(0, w)) ) * QW(w)
}
(psi_Pi_zero <- integrate(integrand, lower = 0, upper = 1)$val)

## [1] 0.1966687
```

(easy algebra reveals that $\Psi(\Pi_0) = 59/300$ indeed).





Figure 1: Evolution of statistical parameter Ψ along fluctuation $\{\Pi_h : h \in H\}$.

2.4 Being smooth, first pass. Luckily, the statistical mapping Ψ is well behaved, or smooth. Here, this colloquial expression refers to the fact that, for each $P \in \mathcal{M}$, if $P_h \rightarrow_h P$ in \mathcal{M} from a direction s when the real parameter $h \rightarrow 0$, then not only $\Psi(P_h) \rightarrow_h \Psi(P)$ (continuity), but also $h^{-1}[\Psi(P_h) - \Psi(P)] \rightarrow_h c$, where the real number c depends on P and s (differentiability).

For instance, let $\Pi_h \in \mathcal{M}$ be the law encoded in `draw_from_another_experiment` with `h` ranging over $[-1, 1]$. We will argue shortly that $\Pi_h \rightarrow_h \Pi_0$ in \mathcal{M} from a direction s when $h \rightarrow 0$. The following chunk of code evaluates and represents $\Psi(\Pi_h)$ for h ranging in a discrete approximation of $[-1, 1]$:

```
approx <- seq(-1, 1, length.out = 1e2)
psi_Pi_h <- sapply(approx, function(t) {
  obs_from_another_experiment <- draw_from_another_experiment(1, h = t)
  integrand <- function(w) {
    Qbar <- attr(obs_from_another_experiment, "Qbar")
    QW <- attr(obs_from_another_experiment, "QW")
    ( Qbar(cbind(1, w)) - Qbar(cbind(0, w)) ) * QW(w)
  }
  integrate(integrand, lower = 0, upper = 1)$val
})
slope_approx <- (psi_Pi_h - psi_Pi_zero) / approx
slope_approx <- slope_approx[min(which(approx > 0))]
ggplot() +
  geom_point(data = data.frame(x = approx, y = psi_Pi_h), aes(x, y),
    color = "#CC6666") +
  geom_segment(aes(x = -1, y = psi_Pi_zero - slope_approx,
    xend = 1, yend = psi_Pi_zero + slope_approx),
    arrow = arrow(length = unit(0.03, "npc")),
    color = "#9999CC") +
  geom_vline(xintercept = 0, color = "#66CC99") +
  geom_hline(yintercept = psi_Pi_zero, color = "#66CC99") +
  labs(x = "h", y = expression(Psi(Pi[h])))
```

The dotted curve represents the function $h \mapsto \Psi(\Pi_h)$. The blue line represents the tangent to the previous curve at $h = 0$, which is indeed differentiable around $h = 0$. It is derived by simple geometric arguments. In the next subsection, we formalize what it means to be smooth for the statistical mapping Ψ . Once the presentation is complete, we will be able to derive a closed-form expression for the slope of the blue curve from the chunk of code where `draw_from_another_experiment` is defined.

2.5  Being smooth, second pass. Let us now describe what it means for statistical mapping Ψ to be smooth at every $P \in \mathcal{M}$. The description necessitates the introduction of fluctuations.

For every direction* $s : \mathcal{O} \rightarrow \mathbb{R}$ such that $s \neq 0^\dagger$, $E_P(s(O)) = 0$ and s bounded by, say, M , for every $h \in H \equiv]-M^{-1}, M^{-1}[$, we can define a law $P_h \in \mathcal{M}$ by setting $P_h \ll P^\ddagger$ and

$$\frac{dP_h}{dP}(O) \equiv 1 + hs(O), \quad (3)$$

that is, P_h has density $(1 + hs)$ with respect to (w.r.t.) P . We call $\{P_h : h \in H\}$ a fluctuation of P in direction s because

$$(i) P_h|_{h=0} = P, \quad (ii) \left. \frac{d}{dh} \log \frac{dP_h}{dP}(O) \right|_{h=0} = s(O). \quad (4)$$

The fluctuation is a one-dimensional parametric submodel of \mathcal{M} .

Statistical mapping Ψ is smooth at every $P \in \mathcal{M}$ because, for each $P \in \mathcal{M}$, there exists a so called efficient influence curve[§] $D^*(P) : \mathcal{O} \rightarrow \mathbb{R}$ such that $E_P(D^*(P)(O)) = 0$ and, for any direction s as above, if $\{P_h : h \in H\}$ is defined as in (3), then the real-valued mapping $h \mapsto \Psi(P_h)$ is differentiable at $h = 0$, with a derivative equal to

$$E_P(D^*(P)(O)s(O)). \quad (5)$$

Interestingly, if a fluctuation $\{P_h : h \in H\}$ satisfies (4) for a direction s such that $s \neq 0$, $E_P(s(O)) = 0$ and $\text{Var}_P(s(O)) < \infty$, then $h \mapsto \Psi(P_h)$ is still differentiable at $h = 0$ with a derivative equal to (5) (beyond fluctuations of the form (3)).

The influence curves $D^*(P)$ convey valuable information about Ψ . For instance, an important result from the theory of inference based on semiparametric models guarantees that if ψ_n is a regular[¶] estimator of $\Psi(P)$ built from n independent observations drawn from P , then the asymptotic variance of the centered and rescaled $\sqrt{n}(\psi_n - \Psi(P))$ cannot be smaller than the variance of the P -specific efficient influence curve, that is,

$$\text{Var}_P(D^*(P)(O)). \quad (6)$$

In this light, an estimator ψ_n of $\Psi(P)$ is said *asymptotically efficient* at P if it is regular at P and such that $\sqrt{n}(\psi_n - \Psi(P))$ converges in law to the centered Gaussian law with variance (6), which is called the Cramér-Rao bound.

2.6 The efficient influence curve. It is not difficult to check (do we give the proof?) that the efficient influence curve $D^*(P)$ of Ψ at $P \in \mathcal{M}$ writes as $D^*(P) \equiv D_1^*(P) + D_2^*(P)$ where $D_1^*(P)$ and $D_2^*(P)$ are given by

*A direction is a measurable function.

†That is, $s(O)$ is not equal to zero P -almost surely.

‡That is, P_h is dominated by P : if an event A satisfies $P(A) = 0$, then necessarily $P_h(A) = 0$ too.

§It is a measurable function.

¶We can view ψ_n as the by product of an algorithm $\hat{\Psi}$ trained on independent observations O_1, \dots, O_n drawn from P . The estimator is regular at P (w.r.t. the maximal tangent space) if, for any direction $s \neq 0$ such that $E_P(s(O)) = 0$ and $\text{Var}_P(s(O)) < \infty$ and fluctuation $\{P_h : h \in H\}$ satisfying (4), the estimator $\psi_{n,1/\sqrt{n}}$ of $\Psi(P_{1/\sqrt{n}})$ obtained by training $\hat{\Psi}$ on independent observations O_1, \dots, O_n drawn from $P_{1/\sqrt{n}}$ is such that $\sqrt{n}(\psi_{n,1/\sqrt{n}} - \Psi(P_{1/\sqrt{n}}))$ converges in law to a limit that does not depend on s .

$$D_1^*(P)(O) \equiv \bar{Q}(1, W) - \bar{Q}(0, W) - \Psi(P),$$

$$D_2^*(P)(O) \equiv \frac{2A - 1}{\ell\bar{G}(A, W)}(Y - \bar{Q}(A, W)),$$

with shorthand notation $\ell\bar{G}(A, W) \equiv A\bar{G}(W) + (1 - A)(1 - \bar{G}(W))$. The following chunk of code enables the computation of the values of the efficient influence curve $D^*(P)$ at observations drawn from P (note that it is necessary to provide the value of $\Psi(P)$, or a numerical approximation thereof, through argument `psi`).

```
eic <- function(obs, psi) {
  Qbar <- attr(obs, "Qbar")
  Gbar <- attr(obs, "Gbar")
  QAW <- Qbar(obs[, c("A", "W")])
  gW <- Gbar(obs[, "W"])
  lgAW <- obs[, "A"] * gW + (1 - obs[, "A"]) * (1 - gW)
  ( Qbar(cbind(1, obs[, "W"])) - Qbar(cbind(0, obs[, "W"])) - psi ) +
    (2 * obs[, "A"] - 1) / lgAW * (obs[, "Y"] - QAW)
}

(eic(five_obs, psi = psi_hat))

## [1] -0.7207719  0.4762960 -0.4398213  0.3962634  0.6449063
(eic(five_obs_from_another_experiment, psi = psi_Pi_zero))

## [1]  0.02107056 -0.00342964  0.10731746  0.07596022  0.05989993
```

2.7 Computing and comparing Cramér-Rao bounds. We can use `eic` to numerically approximate the Cramér-Rao bound at P_0 :

```
obs <- draw_from_experiment(B)
(cramer_rao_hat <- var(eic(obs, psi = psi_hat)))

## [1] 0.3652997
```

and the Cramér-Rao bound at Π_0 :

```
obs_from_another_experiment <- draw_from_another_experiment(B)
(cramer_rao_Pi_zero_hat <- var(eic(obs_from_another_experiment, psi = 59/300)))

## [1] 0.09574321
(ratio <- sqrt(cramer_rao_Pi_zero_hat/cramer_rao_hat))

## [1] 0.5119521
```

We thus discover that of the statistical parameters $\Psi(P_0)$ and $\Psi(\Pi_0)$, the latter is easier to target than the former. Heuristically, for large sample sizes, the narrowest (efficient) confidence intervals for $\Psi(\Pi_0)$ are approximately 0.51 (rounded to two decimal places) smaller than their counterparts for $\Psi(P_0)$.

2.8 Revisiting Section 2.4. It is not difficult either (though a little cumbersome) (do we give the proof? I'd rather not) to verify that $\{\Pi_h : h \in [-1, 1]\}$ is a fluctuation of Π_0 in the direction of σ_0 (in the sense of (3)) given, up to a constant, by

$$\sigma_0(O) \equiv -10\sqrt{W}A \times \beta_0(A, W) \left(\log(1 - Y) + \sum_{k=0}^3 (k + \beta_0(A, W))^{-1} \right) + \text{constant},$$

where $\beta_0(A, W) \equiv \frac{1 - \bar{Q}_{\Pi_0}(A, W)}{\bar{Q}_{\Pi_0}(A, W)}.$

Consequently, the slope of the dotted curve in Figure 1 is equal to

$$E_{\Pi_0}(D^*(\Pi_0)(O)\sigma_0(O)) \quad (7)$$

(since $D^*(\Pi_0)$ is centered under Π_0 , knowing σ_0 up to a constant is not problematic).

Let us check this numerically. In the next chunk of code, we implements direction s with `s_draw_from_another_experiment`, then we numerically approximate (7) (pointwise and with a confidence interval of asymptotic level 95%):

```
s_draw_from_another_experiment <- function(obs) {
  ## preliminary
  Qbar <- attr(obs, "Qbar")
  QAW <- Qbar(obs[, c("A", "W")])
  shape1 <- Arguments$getInteger(attr(obs, "shape1"), c(1, Inf))
  ## computations
  betaAW <- shape1 * (1 - QAW) / QAW
  out <- log(1 - obs[, "Y"])
  for (int in 1:shape1) {
    out <- out + 1/(int - 1 + betaAW)
  }
  out <- - out * shape1 * (1 - QAW) / QAW * 10 * sqrt(obs[, "W"]) * obs[, "A"]
  ## no need to center given how we will use it
  return(out)
}

vars <- eic(obs_from_another_experiment, psi = 59/300) *
  s_draw_from_another_experiment(obs_from_another_experiment)
sd_hat <- sd(vars)
(slope_hat <- mean(vars))

## [1] 1.358524

(slope_CI <- slope_hat + c(-1, 1) * qnorm(1 - alpha / 2) * sd_hat / sqrt(B))

## [1] 1.353257 1.363791
```

Equal to 1.349 (rounded to three decimal places), the first numerical approximation `slope_approx` is not too off.

2.9 Double-robustness The efficient influence curve $D^*(P)$ at $P \in \mathcal{M}$ enjoys another remarkable property: it is double-robust. Specifically, for all $P' \in \mathcal{M}$, it holds that

$$\Psi(P') - \Psi(P) = -E_P(D^*(P')(O)) + \text{Rem}_P(\bar{Q}', \bar{G}') \quad (8)$$

where the so called remainder term $\text{Rem}_P(\bar{Q}', \bar{G}')$ satisfies^{||}

^{||}For any (measurable) $f : \mathcal{O} \rightarrow \mathbb{R}$, we denote $\|f\|_P = E_P(f(O)^2)^{1/2}.$

$$\text{Rem}_P(\bar{Q}', \bar{G}')^2 \leq \|\bar{Q}' - \bar{Q}\|_P^2 \times \|(\bar{G}' - \bar{G})/\ell\bar{G}'\|_P^2. \quad (9)$$

In particular, if

$$E_P(D^*(P')(O)) = 0, \quad (10)$$

and *either* $\bar{Q}' = \bar{Q}$ or $\bar{G}' = \bar{G}$, then $\text{Rem}_P(\bar{Q}', \bar{G}') = 0$ hence $\Psi(P') = \Psi(P)$. In words, if P' solves the so called P -specific efficient influence curve equation (10) and if, in addition, P' has the same \bar{Q} -component or \bar{G} -component as P , then $\Psi(P') = \Psi(P)$ no matter how P' may differ from P otherwise. This property is useful to build consistent estimators of $\Psi(P)$.

However, there is much more to double-robustness than the above straightforward implication. Indeed, 8 is useful to build a consistent estimator of $\Psi(P)$ that, in addition, satisfies a central limit theorem and thus lends itself to the construction of confidence intervals.

Let $P_n^0 \in \mathcal{M}$ be an element of model \mathcal{M} of which the choice is data-driven, based on observing n independent draws from P . Equality 8 reveals that the statistical behavior of the corresponding *substitution* estimator $\psi_n^0 \equiv \Psi(P_n^0)$ is easier to analyze when the remainder term $\text{Rem}_P(\bar{Q}_n^0, \bar{G}_n^0)$ goes to zero at a fast (relative to n) enough rate. In light of 9, this happens if the features \bar{Q}_n^0 and \bar{G}_n^0 of P_n^0 converge to their counterparts under P at rates of which *the product* is fast enough.

2.10 Targeted inference. blabla

```
Gbar <- attr(obs, "Gbar")

psi_hat_ab <- obs %>% as_tibble() %>% mutate(id = 1:nrow(obs) %>% 1e3) %>%
  mutate(lgAW = A * Gbar(W) + (1 - A) * (1 - Gbar(W))) %>% group_by(id) %>%
  summarize(est_a = mean(Y[A==1]) - mean(Y[A==0]),
            est_b = mean(Y * (2 * A - 1) / lgAW),
            std_b = sd(Y * (2 * A - 1) / lgAW),
            clt_b = sqrt(n()) * (est_b - psi_hat) / std_b)
std_a <- sd(psi_hat_ab$est_a)
psi_hat_ab <- psi_hat_ab %>%
  mutate(std_a = std_a,
         clt_a = (est_a - psi_hat) / std_a) %>%
  gather(key, value, -id) %>%
  extract(key, c("what", "type"), "([^-]+)_([ab])") %>%
  spread(what, value)

bias <- psi_hat_ab %>% group_by(type) %>% summarise(bias = mean(clt))

ggplot(psi_hat_ab, aes(clt, fill = type, colour = type)) +
  geom_density(alpha = 0.1) +
  geom_vline(aes(xintercept = bias, colour = type), bias) +
  labs(x = expression(paste(sqrt(n)*(psi[n]^{list(a, b)} - psi[0]))))
```

blabla




Figure 2: Of two estimators of ψ_0 , one of them misconceived, the other assuming that \bar{G}_0 is known.