






A guided tour in targeted learning territory

David Benkeser, Antoine Chambaz, Nima Hejazi

08/20/2018

Contents

1	Introduction	1
2	A simulation study	2
2.1	Reproducible experiment as a law.	2
2.2	 Visualizing infinite-dimensional features of the experiment	4
2.3	The parameter of interest, first pass.	5
2.4	 Difference in covariate-adjusted quantile rewards, first pass.	7
2.5	The parameter of interest, second pass.	7
2.6	 Difference in covariate-adjusted quantile rewards, second pass.	9
2.7	Being smooth, first pass.	10
2.8	 Being smooth, second pass.	11
2.9	The efficient influence curve.	12
2.10	Computing and comparing Cramér-Rao bounds.	12
2.11	Revisiting Section 2.7.	13
2.12	Double-robustness	14
2.13	Inference assuming \bar{G}_0 known, or not, first pass.	14
2.14	Inference assuming \bar{G}_0 known, or not, second pass.	17
2.15	 Exercises.	19
2.16	Inference based on the estimation of \bar{Q}_0	21
2.17	One-step estimation.	25
2.18	Targeted inference.	25

1 Introduction

This is a very first draft of our article. The current *tentative* title is "A guided tour in targeted learning territory".

Explain our objectives and how we will meet them. Explain that the symbol  indicates more delicate material.

Use sectioning a lot to ease cross-referencing.

Do we include exercises? I propose we do, and to flag the corresponding subsections with symbol .

```

set.seed(54321) ## because reproducibility matters...
suppressMessages(library(R.utils)) ## make sure it is installed
suppressMessages(library(tidyverse)) ## make sure it is installed
suppressMessages(library(caret)) ## make sure it is installed
expit <- plogis
logit <- qlogis

```

Function `expit` implements the link function $\text{expit} : \mathbb{R} \rightarrow]0, 1[$ given by $\text{expit}(x) \equiv (1 + e^{-x})^{-1}$. Function `logit` implements its inverse function $\text{logit} :]0, 1[\rightarrow \mathbb{R}$ given by $\text{logit}(p) \equiv \log[p/(1 - p)]$.

2 A simulation study

blabla

2.1 Reproducible experiment as a law. We are interested in a reproducible experiment. The generic summary of how one realization of the experiment unfolds, our observation, is called O . We view O as a random variable drawn from what we call the law P_0 of the experiment. The law P_0 is viewed as an element of what we call the model. Denoted by \mathcal{M} , the model is the collection of *all* laws from which O can be drawn and that meet some constraints. The constraints translate the knowledge we have about the experiment. The more we know about the experiment, the smaller is \mathcal{M} . In all our examples, model \mathcal{M} will put very few restrictions on the candidate laws.

Consider the following chunk of code:

```

draw_from_experiment <- function(n, ideal = FALSE) {
  ## preliminary
  n <- Arguments$getInteger(n, c(1, Inf))
  ideal <- Arguments$getLogical(ideal)
  ## ## 'Gbar' and 'Qbar' factors
  Gbar <- function(W) {
    expit(1 + 2 * W - 4 * sqrt(abs((W - 5/12))))
  }
  Qbar <- function(AW) {
    A <- AW[, 1]
    W <- AW[, 2]
    ## A * (cos((1 + W) * pi / 4) + (1/3 <= W & W <= 1/2) / 5) +
    ## (1 - A) * (sin(4 * W^2 * pi) / 4 + 1/2)
    A * (cos((-1/2 + W) * pi) * 2/5 + 1/5 + (1/3 <= W & W <= 1/2) / 5 +
      (W >= 3/4) * (W - 3/4) * 2) +
    (1 - A) * (sin(4 * W^2 * pi) / 4 + 1/2)
  }
  ## sampling
  ## ## context
  mixture_weights <- c(1/10, 9/10, 0)
  mins <- c(0, 11/30, 0)
  maxs <- c(1, 14/30, 1)
  latent <- findInterval(runif(n), cumsum(mixture_weights)) + 1
  W <- runif(n, min = mins[latent], max = maxs[latent])
  ## ## counterfactual rewards
  zeroW <- cbind(A = 0, W)

```

```

oneW <- cbind(A = 1, W)
Qbar.zeroW <- Qbar(zeroW)
Qbar.oneW <- Qbar(oneW)
Yzero <- rbeta(n, shape1 = 2, shape2 = 2 * (1 - Qbar.zeroW) / Qbar.zeroW)
Yone <- rbeta(n, shape1 = 3, shape2 = 3 * (1 - Qbar.oneW) / Qbar.oneW)
## ## action undertaken
A <- rbinom(n, size = 1, prob = Gbar(W))
## ## actual reward
Y <- A * Yone + (1 - A) * Yzero
## ## observation
if (ideal) {
  obs <- cbind(W = W, Yzero = Yzero, Yone = Yone, A = A, Y = Y)
} else {
  obs <- cbind(W = W, A = A, Y = Y)
}
attr(obs, "Gbar") <- Gbar
attr(obs, "Qbar") <- Qbar
attr(obs, "QW") <- function(W) {
  out <- sapply(1:length(mixture_weights),
    function(ii){
      mixture_weights[ii] *
      dunif(W, min = mins[ii], max = maxs[ii])
    })
  return(rowSums(out))
}
attr(obs, "qY") <- function(AW, Y, Qbar){
  A <- AW[, 1]
  W <- AW[, 2]
  Qbar.AW <- do.call(Qbar, list(AW)) # is call to 'do.call' necessary?
  shape1 <- ifelse(A == 0, 2, 3)
  dbeta(Y, shape1 = shape1, shape2 = shape1 * (1 - Qbar.AW) / Qbar.AW)
}
##
return(obs)
}

```

We can interpret `draw_from_experiment` as a law P_0 since we can use the function to sample observations from a common law. It is even a little more than that, because we can tweak the experiment, by setting its `ideal` argument to `TRUE`, in order to get what appear as intermediary (counterfactual) variables in the regular experiment. The next chunk of code runs the (regular) experiment five times independently and outputs the resulting observations:

```

(five_obs <- draw_from_experiment(5))

```

```

##           W A           Y
## [1,] 0.4533028 0 0.8979460
## [2,] 0.3716077 0 0.9905312
## [3,] 0.3875802 0 0.8080567
## [4,] 0.4008279 1 0.9954100
## [5,] 0.4038325 0 0.9772926
## attr(,"Gbar")
## function (W)
## {

```

```

##      expit(1 + 2 * W - 4 * sqrt(abs((W - 5/12))))
## }
## <bytecode: 0x3e245e58>
## <environment: 0x418fc550>
## attr("Qbar")
## function (AW)
## {
##     A <- AW[, 1]
##     W <- AW[, 2]
##     A * (cos((-1/2 + W) * pi) * 2/5 + 1/5 + (1/3 <= W & W <=
##         1/2)/5 + (W >= 3/4) * (W - 3/4) * 2) + (1 - A) * (sin(4 *
##         W^2 * pi)/4 + 1/2)
## }
## <bytecode: 0x48851848>
## <environment: 0x418fc550>
## attr("QW")
## function (W)
## {
##     out <- sapply(1:length(mixture_weights), function(ii) {
##         mixture_weights[ii] * dunif(W, min = mins[ii], max = maxs[ii])
##     })
##     return(rowSums(out))
## }
## <bytecode: 0x3e132470>
## <environment: 0x418fc550>
## attr("qY")
## function (AW, Y, Qbar)
## {
##     A <- AW[, 1]
##     W <- AW[, 2]
##     Qbar.AW <- do.call(Qbar, list(AW))
##     shape1 <- ifelse(A == 0, 2, 3)
##     dbeta(Y, shape1 = shape1, shape2 = shape1 * (1 - Qbar.AW)/Qbar.AW)
## }
## <bytecode: 0x3e89dcb8>
## <environment: 0x418fc550>

```

We can view the `attributes` of object `five_obs` because, in this section, we act as oracles, *i.e.*, we know completely the nature of the experiment. In particular, we have included several features of P_0 that play an important role in our developments. The attribute `QW` describes the density of W , of which the law $Q_{0,W}$ is a mixture of the uniform laws over $[0, 1]$ (weight $1/10$) and $[11/30, 14/30]$ (weight $9/10$).^{*} The attribute `Gbar` describes the conditional probability of action $A = 1$ given W . For each $a \in \{0, 1\}$, we denote $\bar{G}_0(W) \equiv \Pr_{P_0}(A = 1|W)$ and $\ell\bar{G}_0(a, W) \equiv \Pr_{P_0}(A = a|W)$. The attribute `qY` describes the conditional density of Y given A and W . For each $y \in]0, 1[$, we denote by $q_{0,Y}(y, A, W)$ the conditional density evaluated at y of Y given A and W . Similarly, the attribute `Qbar` describes the conditional mean of Y given A and W , and we denote $\bar{Q}_0(A, W) = E_{P_0}(Y|A, W)$ the conditional mean of Y given A and W .

For consistency, I removed the $A = a$ and $W = w$.

2.2 Visualizing infinite-dimensional features of the experiment


^{*}We fine-tuned (or tweaked, or something else?) the marginal law of W to make it easier later on to drive home important messages. Specifically, ... (do we explain what happens?)

1. Run the following chunk of code. It visualizes the conditional mean \bar{Q}_0 .

```
Gbar <- attr(five_obs, "Gbar")
Qbar <- attr(five_obs, "Qbar")
QW <- attr(five_obs, "QW")

features <- tibble(w = seq(0, 1, length.out = 1e3)) %>%
  mutate(Qw = QW(w),
         Gw = Gbar(w),
         Q1w = Qbar(cbind(A = 1, W = w)),
         Q0w = Qbar(cbind(A = 0, W = w)),
         blip_Qw = Q1w - Q0w)

features %>% select(-Qw, -Gw) %>%
  rename("Q(1,.)" = Q1w,
        "Q(0,.)" = Q0w,
        "Q(1,.) - Q(0,.)" = blip_Qw) %>%
  gather("f", "value", -w) %>%
  ggplot() +
  geom_line(aes(x = w, y = value, color = f), size = 1) +
  labs(y = "f(w)", title = bquote("Visualizing" ~ bar(Q)[0])) +
  ylim(NA, 1)
```



2. Adapt the above chunk of code to visualize the marginal density $Q_{0,W}$ and conditional probability \bar{G}_0 .

2.3 The parameter of interest, first pass. It happens that we especially care for a finite-dimensional feature of P_0 that we denote by ψ_0 . Its definition involves two of the aforementioned infinite-dimensional features:

$$\begin{aligned}\psi_0 &\equiv \int (\bar{Q}_0(1, w) - \bar{Q}_0(0, w)) dQ_{0,W}(w) \\ &= E_{P_0} (\bar{Q}_0(1, W) - \bar{Q}_0(0, W)).\end{aligned}\tag{1}$$

Acting as oracles, we can compute explicitly the numerical value of ψ_0 .

```
integrand <- function(w) {
  Qbar <- attr(five_obs, "Qbar")
```

```

QW <- attr(five_obs, "QW")
( Qbar(cbind(1, w)) - Qbar(cbind(0, w)) ) * QW(w)
}
(psi_zero <- integrate(integrand, lower = 0, upper = 1)$val)

```

```
## [1] 0.08317711
```

Our interest in ψ_0 is of causal nature. Taking a closer look at `drawFromExperiment` reveals indeed that the random making of an observation O drawn from P_0 can be summarized by the following causal graph and nonparametric system of structural equations:

```
## plot the causal diagram
```

and, for some deterministic functions f_w, f_a, f_y and independent sources of randomness U_w, U_a, U_y ,

1. sample the context where the rest of the experiment will take place, $W = f_w(U_w)$;
2. sample the two counterfactual rewards of the two actions that can be undertaken, $Y_0 = f_y(0, W, U_y)$ and $Y_1 = f_y(1, W, U_y)$;
3. sample which action is carried out in the given context, $A = f_a(W, U_a)$;
4. define the corresponding reward, $Y = AY_1 + (1 - A)Y_0$;
5. summarize the course of the experiment with the observation $O = (W, A, Y)$, thus concealing Y_0 and Y_1 .

The above description of the experiment `draw_from_experiment` is useful to reinforce what it means to run the “ideal” experiment by setting argument `ideal` to `TRUE` in a call to `draw_from_experiment`. Doing so triggers a modification of the nature of the experiment, enforcing that the counterfactual rewards Y_0 and Y_1 be part of the summary of the experiment eventually. In light of the above enumeration, $\mathbb{O} \equiv (W, Y_0, Y_1, A, Y)$ is output, as opposed to its summary measure O . This defines another experiment and its law, that we denote \mathbb{P}_0 .

It is well known (do we give the proof or refer to other articles?) that

$$\psi_0 = \mathbb{E}_{\mathbb{P}_0}(Y_1 - Y_0) = \mathbb{E}_{\mathbb{P}_0}(Y_1) - \mathbb{E}_{\mathbb{P}_0}(Y_0). \quad (2)$$

Thus, ψ_0 describes the average difference in of the two counterfactual rewards. In other words, ψ_0 quantifies the difference in average of the reward one would get in a world where one would always enforce action $a = 1$ with the reward one would get in a world where one would always enforce action $a = 0$. This said, it is worth emphasizing that ψ_0 is a well defined parameter beyond its causal interpretation.

To conclude this subsection, we take advantage of our position as oracles to sample observations from the ideal experiment. We call `draw_from_experiment` with its argument `ideal` set to `TRUE` in order to numerically approximate ψ_0 . By the law of large numbers, the following chunk of code approximates ψ_0 and shows its approximate value:

```

B <- 1e6 ## Antoine: 1e6 eventually
ideal_obs <- draw_from_experiment(B, ideal = TRUE)
(psi_approx <- mean(ideal_obs[, "Yone"] - ideal_obs[, "Yzero"]))


```

```
## [1] 0.08293116
```

In fact, the central limit theorem and Slutsky's lemma allow us to build a confidence interval with asymptotic level 95% for ψ_0 :

```
sd_approx <- sd(ideal_obs[, "Yone"] - ideal_obs[, "Yzero"])
alpha <- 0.05
(psi_approx_CI <- psi_approx + c(-1, 1) * qnorm(1 - alpha / 2) * sd_approx / sqrt(B))
```

```
## [1] 0.08232477 0.08353754
```

2.4  Difference in covariate-adjusted quantile rewards, first pass. The problems come within the scope of Sections 2.3.

As discussed above, parameter ψ_0 (2) is the difference in average rewards if we enforce action $a = 1$ rather than $a = 0$. An alternative way to describe the rewards under different actions involves quantiles as opposed to averages.

Let $Q_{0,Y}(y, a, w) = \int_0^y q_{0,Y}(u, a, w) du$ be the conditional cumulative distribution of reward Y given $A = a$ and $W = w$, evaluated at $y \in]0, 1[$, that is implied by P_0 . For each action $a \in \{0, 1\}$ and $c \in]0, 1[$, introduce

$$\gamma_{0,a,c} \equiv \inf \left\{ y \in]0, 1[: \int Q_{0,Y}(y, a, w) dQ_{0,W}(w) \geq c \right\}. \quad (3)$$


It is not difficult to check that [do we give the proof or refer to other articles?](#)

$$\gamma_{0,a,c} = \inf \left\{ y \in]0, 1[: \Pr_{\mathbb{P}_0}(Y_a \leq y) \geq c \right\}.$$

Thus, $\gamma_{0,a,c}$ can be interpreted as a covariate-adjusted c -th quantile reward when action a is enforced. The difference

$$\delta_{0,c} \equiv \gamma_{0,1,c} - \gamma_{0,0,c}$$

is the c -th quantile counterpart to parameter ψ_0 (2).

1.  Compute the numerical value of $\gamma_{0,a,c}$ for each $(a, c) \in \{0, 1\} \times \{1/4, 1/2, 3/4\}$ using the appropriate attributes of `five_obs`. Based on these results, report the numerical value of $\delta_{0,c}$ for each $c \in \{1/4, 1/2, 3/4\}$.
2. Approximate the numerical values of $\gamma_{0,a,c}$ for each $(a, c) \in \{0, 1\} \times \{1/4, 1/2, 3/4\}$ by drawing a large sample from the “ideal” data experiment and using empirical quantile estimates. Deduce from these results a numerical approximation to $\delta_{0,c}$ for $c \in \{1/4, 1/2, 3/4\}$. Confirm that your results closely match those obtained in the previous problem.

2.5 The parameter of interest, second pass. Suppose we know beforehand that O drawn from P_0 takes its values in $\mathcal{O} \equiv [0, 1] \times \{0, 1\} \times [0, 1]$ and that $\bar{G}(W) = P_0(A = 1|W)$ is bounded away from zero and one $Q_{0,W}$ -almost surely (this is the case indeed). Then we can define model \mathcal{M} as the set of all laws P on \mathcal{O}

such that $\bar{G}(W) \equiv P(A = 1|W)$ is bounded away from zero and one Q_W -almost surely, where Q_W is the marginal law of W under P .

Let us also define generically \bar{Q} as

$$\bar{Q}(A, W) \equiv E_P(Y|A, W).$$

Central to our approach is viewing ψ_0 as the value at P_0 of the statistical mapping Ψ from \mathcal{M} to $[0, 1]$ characterized by

$$\begin{aligned}\Psi(P) &\equiv \int (\bar{Q}(1, w) - \bar{Q}(0, w)) dQ_W(w) \\ &= E_P(\bar{Q}(1, W) - \bar{Q}(0, W)),\end{aligned}$$

a clear extension of (1). For instance, although the law $\Pi_0 \in \mathcal{M}$ encoded by default (*i.e.*, with $h=0$) in `drawFromAnotherExperiment` defined below differs starkly from P_0 ,

```
draw_from_another_experiment <- function(n, h = 0) {
  ## preliminary
  n <- Arguments$getInteger(n, c(1, Inf))
  h <- Arguments$getNumeric(h)
  ## ## 'Gbar' and 'Qbar' factors
  Gbar <- function(W) {
    sin((1 + W) * pi / 6)
  }
  Qbar <- function(AW, hh = h) {
    A <- AW[, 1]
    W <- AW[, 2]
    expit( logit( A * W + (1 - A) * W^2 ) +
            hh * 10 * sqrt(W) * A )
  }
  ## sampling
  ## ## context
  W <- runif(n, min = 1/10, max = 9/10)
  ## ## action undertaken
  A <- rbinom(n, size = 1, prob = Gbar(W))
  ## ## reward
  shape1 <- 4
  QAW <- Qbar(cbind(A, W))
  Y <- rbeta(n, shape1 = shape1, shape2 = shape1 * (1 - QAW) / QAW)
  ## ## observation
  obs <- cbind(W = W, A = A, Y = Y)
  attr(obs, "Gbar") <- Gbar
  attr(obs, "Qbar") <- Qbar
  attr(obs, "QW") <- function(x){dunif(x, min = 1/10, max = 9/10)}
  attr(obs, "shape1") <- shape1
  attr(obs, "qY") <- function(AW, Y, Qbar, shape1){
    A <- AW[,1]; W <- AW[,2]
    Qbar.AW <- do.call(Qbar, list(AW))
    dbeta(Y, shape1 = shape1, shape2 = shape1 * (1 - Qbar.AW) / Qbar.AW)
```



```

}
##
return(obs)
}

```

the parameter $\Psi(\Pi_0)$ is well defined, and numerically approximated by `psi_Pi_zero` as follows.


```

five_obs_from_another_experiment <- draw_from_another_experiment(5)
another_integrand <- function(w) {
  Qbar <- attr(five_obs_from_another_experiment, "Qbar")
  QW <- attr(five_obs_from_another_experiment, "QW")
  ( Qbar(cbind(1, w)) - Qbar(cbind(0, w)) ) * QW(w)
}
(psi_Pi_zero <- integrate(another_integrand, lower = 0, upper = 1)$val)

```

```
## [1] 0.1966687
```

Straightforward algebra confirms that indeed $\Psi(\Pi_0) = 59/300$.

 **2.6 Difference in covariate-adjusted quantile rewards, second pass.** We continue with the exercise from Section 2.4. The problems come within the scope of Section 2.3.

As above, we define $q_Y(y, a, w)$ to be the (A, W) -conditional density of Y given $A = a$ and $W = w$, evaluated at y , that is implied by a generic $P \in \mathcal{M}$. Similarly, we use Q_Y to denote the corresponding cumulative distribution function. The covariate-adjusted c -th quantile reward for action $a \in \{0, 1\}$ may be viewed as a mapping $\Gamma_{a,c}$ from \mathcal{M} to $[0, 1]$ characterized by

$$\Gamma_{a,c}(P) = \inf \left\{ y \in]0, 1[: \int Q_Y(y, a, w) dQ_W(w) \geq c \right\}.$$

The difference in c -th quantile rewards may similarly be viewed as a mapping Δ_c from \mathcal{M} to $[0, 1]$, characterized by $\Delta_c(P) \equiv \Gamma_{1,c}(P) - \Gamma_{0,c}(P)$.

1. Compute the numerical value of $\Gamma_{a,c}(\Pi_0)$ for $(a, c) \in \{0, 1\} \times \{1/4, 1/2, 3/4\}$ using the appropriate attributes of `five_obs_from_another_experiment`. Based on these results, report the numerical value of $\Delta_c(\Pi_0)$ for each $c \in \{1/4, 1/2, 3/4\}$.
2. Approximate the value of $\Gamma_{0,a,c}(\Pi_0)$ for $(a, c) \in \{0, 1\} \times \{1/4, 1/2, 3/4\}$ by drawing a large sample from the “ideal” data experiment and using empirical quantile estimates. Deduce from these results a numerical approximation to $\Delta_{0,c}(\Pi_0)$ for each $c \in \{1/4, 1/2, 3/4\}$. Confirm that your results closely match those obtained in the previous problem.
3. Building upon the code you wrote to solve the previous problem, construct a confidence interval with asymptotic level 95% for $\Delta_{0,c}(\Pi_0)$, with $c \in \{1/4, 1/2, 3/4\}$.[†]

[†]Let X_1, \dots, X_n be independently drawn from a continuous distribution function F . Set $p \in]0, 1[$ and, assuming that n is large, find $k \geq 1$ and $l \geq 1$ such that $k/n \approx p - \Phi^{-1}(1 - \alpha)\sqrt{p(1-p)/n}$ and $l/n \approx p + \Phi^{-1}(1 - \alpha)\sqrt{p(1-p)/n}$, where Φ is the standard normal distribution function. Then $[X_{(k)}, X_{(l)}]$ is a confidence interval for $F^{-1}(p)$ with asymptotic level $1 - 2\alpha$.





Figure 1: Evolution of statistical parameter Ψ along fluctuation $\{\Pi_h : h \in H\}$.

2.7 Being smooth, first pass. Luckily, the statistical mapping Ψ is well behaved, or smooth. Here, this colloquial expression refers to the fact that, for each $P \in \mathcal{M}$, if $P_h \rightarrow_h P$ in \mathcal{M} from a direction s when the real parameter $h \rightarrow 0$, then not only $\Psi(P_h) \rightarrow_h \Psi(P)$ (continuity), but also $h^{-1}[\Psi(P_h) - \Psi(P)] \rightarrow_h c$, where the real number c depends on P and s (differentiability).

For instance, let $\Pi_h \in \mathcal{M}$ be the law encoded in `draw_from_another_experiment` with `h` ranging over $[-1, 1]$. We will argue shortly that $\Pi_h \rightarrow_h \Pi_0$ in \mathcal{M} from a direction s when $h \rightarrow 0$. The following chunk of code evaluates and represents $\Psi(\Pi_h)$ for h ranging in a discrete approximation of $[-1, 1]$:

```
approx <- seq(-1, 1, length.out = 1e2)
psi_Pi_h <- sapply(approx, function(t) {
  obs_from_another_experiment <- draw_from_another_experiment(1, h = t)
  integrand <- function(w) {
    Qbar <- attr(obs_from_another_experiment, "Qbar")
    QW <- attr(obs_from_another_experiment, "QW")
    ( Qbar(cbind(1, w)) - Qbar(cbind(0, w)) ) * QW(w)
  }
  integrate(integrand, lower = 0, upper = 1)$val
})
slope_approx <- (psi_Pi_h - psi_Pi_zero) / approx
slope_approx <- slope_approx[min(which(approx > 0))]
ggplot() +
  geom_point(data = data.frame(x = approx, y = psi_Pi_h), aes(x, y),
    color = "#CC6666") +
  geom_segment(aes(x = -1, y = psi_Pi_zero - slope_approx,
    xend = 1, yend = psi_Pi_zero + slope_approx),
    arrow = arrow(length = unit(0.03, "npc")),
    color = "#9999CC") +
  geom_vline(xintercept = 0, color = "#66CC99") +
  geom_hline(yintercept = psi_Pi_zero, color = "#66CC99") +
  labs(x = "h", y = expression(Psi(Pi[h])))
```

The dotted curve represents the function $h \mapsto \Psi(\Pi_h)$. The blue line represents the tangent to the previous curve at $h = 0$, which is indeed differentiable around $h = 0$. It is derived by simple geometric arguments. In the next subsection, we formalize what it means to be smooth for the statistical mapping Ψ . Once the presentation is complete, we will be able to derive a closed-form expression for the slope of the blue curve from the chunk of code where `draw_from_another_experiment` is defined.

2.8  Being smooth, second pass. Let us now describe what it means for statistical mapping Ψ to be smooth at every $P \in \mathcal{M}$. The description necessitates the introduction of fluctuations.

For every direction[†] $s : \mathcal{O} \rightarrow \mathbb{R}$ such that $s \neq 0^{\S}$, $E_P(s(O)) = 0$ and s bounded by, say, M , for every $h \in H \equiv]-M^{-1}, M^{-1}[$, we can define a law $P_h \in \mathcal{M}$ by setting $P_h \ll P^{\P}$ and

$$\frac{dP_h}{dP}(O) \equiv 1 + hs(O), \quad (4)$$

that is, P_h has density $(1 + hs)$ with respect to (w.r.t.) P . We call $\{P_h : h \in H\}$ a fluctuation of P in direction s because

$$(i) P_h|_{h=0} = P, \quad (ii) \left. \frac{d}{dh} \log \frac{dP_h}{dP}(O) \right|_{h=0} = s(O). \quad (5)$$

The fluctuation is a one-dimensional parametric submodel of \mathcal{M} .

Statistical mapping Ψ is smooth at every $P \in \mathcal{M}$ because, for each $P \in \mathcal{M}$, there exists a so called efficient influence curve^{||} $D^*(P) : \mathcal{O} \rightarrow \mathbb{R}$ such that $E_P(D^*(P)(O)) = 0$ and, for any direction s as above, if $\{P_h : h \in H\}$ is defined as in (4), then the real-valued mapping $h \mapsto \Psi(P_h)$ is differentiable at $h = 0$, with a derivative equal to

$$E_P(D^*(P)(O)s(O)). \quad (6)$$

Interestingly, if a fluctuation $\{P_h : h \in H\}$ satisfies (5) for a direction s such that $s \neq 0$, $E_P(s(O)) = 0$ and $\text{Var}_P(s(O)) < \infty$, then $h \mapsto \Psi(P_h)$ is still differentiable at $h = 0$ with a derivative equal to (6) (beyond fluctuations of the form (4)).

The influence curves $D^*(P)$ convey valuable information about Ψ . For instance, an important result from the theory of inference based on semiparametric models guarantees that if ψ_n is a regular^{**} estimator of $\Psi(P)$ built from n independent observations drawn from P , then the asymptotic variance of the centered and rescaled $\sqrt{n}(\psi_n - \Psi(P))$ cannot be smaller than the variance of the P -specific efficient influence curve, that is,

$$\text{Var}_P(D^*(P)(O)). \quad (7)$$

In this light, an estimator ψ_n of $\Psi(P)$ is said *asymptotically efficient* at P if it is regular at P and such that $\sqrt{n}(\psi_n - \Psi(P))$ converges in law to the centered Gaussian law with variance (7), which is called the Cramér-Rao bound.

[†]A direction is a measurable function.

[§]That is, $s(O)$ is not equal to zero P -almost surely.

[¶]That is, P_h is dominated by P : if an event A satisfies $P(A) = 0$, then necessarily $P_h(A) = 0$ too.

^{||}It is a measurable function.

^{**}We can view ψ_n as the by product of an algorithm $\hat{\Psi}$ trained on independent observations O_1, \dots, O_n drawn from P . The estimator is regular at P (w.r.t. the maximal tangent space) if, for any direction $s \neq 0$ such that $E_P(s(O)) = 0$ and $\text{Var}_P(s(O)) < \infty$ and fluctuation $\{P_h : h \in H\}$ satisfying (5), the estimator $\psi_{n,1/\sqrt{n}}$ of $\Psi(P_{1/\sqrt{n}})$ obtained by training $\hat{\Psi}$ on independent observations O_1, \dots, O_n drawn from $P_{1/\sqrt{n}}$ is such that $\sqrt{n}(\psi_{n,1/\sqrt{n}} - \Psi(P_{1/\sqrt{n}}))$ converges in law to a limit that does not depend on s .

2.9 The efficient influence curve. It is not difficult to check (do we give the proof?) that the efficient influence curve $D^*(P)$ of Ψ at $P \in \mathcal{M}$ writes as $D^*(P) \equiv D_1^*(P) + D_2^*(P)$ where $D_1^*(P)$ and $D_2^*(P)$ are given by

$$D_1^*(P)(O) \equiv \bar{Q}(1, W) - \bar{Q}(0, W) - \Psi(P),$$

$$D_2^*(P)(O) \equiv \frac{2A - 1}{\ell\bar{G}(A, W)}(Y - \bar{Q}(A, W)),$$

with shorthand notation $\ell\bar{G}(A, W) \equiv A\bar{G}(W) + (1 - A)(1 - \bar{G}(W))$. The following chunk of code enables the computation of the values of the efficient influence curve $D^*(P)$ at observations drawn from P (note that it is necessary to provide the value of $\Psi(P)$, or a numerical approximation thereof, through argument `psi`).

```
eic <- function(obs, psi) {
  Qbar <- attr(obs, "Qbar")
  Gbar <- attr(obs, "Gbar")
  QAW <- Qbar(obs[, c("A", "W")])
  GW <- Gbar(obs[, "W"])
  lGAW <- obs[, "A"] * GW + (1 - obs[, "A"]) * (1 - GW)
  ( Qbar(cbind(1, obs[, "W"])) - Qbar(cbind(0, obs[, "W"])) - psi ) +
    (2 * obs[, "A"] - 1) / lGAW * (obs[, "Y"] - QAW)
}

(eic(five_obs, psi = psi_approx))
```

```
## [1] -1.0154401 -0.9019670 -0.3260029 0.2457730 -1.2678877
```

```
(eic(five_obs_from_another_experiment, psi = psi_Pi_zero))
```

```
## [1] -1.14434889 -0.04617269 0.15717587 0.38152278 0.09369052
```

2.10 Computing and comparing Cramér-Rao bounds. We can use `eic` to numerically approximate the Cramér-Rao bound at P_0 :

```
obs <- draw_from_experiment(B)
(cramer_rao_hat <- var(eic(obs, psi = psi_approx)))
```

```
## [1] 0.2874799
```

and the Cramér-Rao bound at Π_0 :

```
obs_from_another_experiment <- draw_from_another_experiment(B)
(cramer_rao_Pi_zero_hat <- var(eic(obs_from_another_experiment, psi = 59/300)))
```

```
## [1] 0.09476356
```

```
(ratio <- sqrt(cramer_rao_Pi_zero_hat/cramer_rao_hat))
```

```
## [1] 0.5741389
```

We thus discover that of the statistical parameters $\Psi(P_0)$ and $\Psi(\Pi_0)$, the latter is easier to target than the former. Heuristically, for large sample sizes, the narrowest (efficient) confidence intervals for $\Psi(\Pi_0)$ are approximately 0.57 (rounded to two decimal places) smaller than their counterparts for $\Psi(P_0)$.

2.11 Revisiting Section 2.7. It is not difficult either (though a little cumbersome) (do we give the proof? I'd rather not) to verify that $\{\Pi_h : h \in [-1, 1]\}$ is a fluctuation of Π_0 in the direction of σ_0 (in the sense of (4)) given, up to a constant, by

$$\sigma_0(O) \equiv -10\sqrt{W}A \times \beta_0(A, W) \left(\log(1 - Y) + \sum_{k=0}^3 (k + \beta_0(A, W))^{-1} \right) + \text{constant},$$

where $\beta_0(A, W) \equiv \frac{1 - \bar{Q}_{\Pi_0}(A, W)}{\bar{Q}_{\Pi_0}(A, W)}.$

Consequently, the slope of the dotted curve in Figure 1 is equal to

$$E_{\Pi_0}(D^*(\Pi_0)(O)\sigma_0(O)) \quad (8)$$

(since $D^*(\Pi_0)$ is centered under Π_0 , knowing σ_0 up to a constant is not problematic).

Let us check this numerically. In the next chunk of code, we implement direction σ_0 with `sigma0_draw_from_another_experiment`, then we numerically approximate (8) (pointwise and with a confidence interval of asymptotic level 95%):

```
sigma0_draw_from_another_experiment <- function(obs) {
  ## preliminary
  Qbar <- attr(obs, "Qbar")
  QAW <- Qbar(obs[, c("A", "W")])
  shape1 <- Arguments$getInteger(attr(obs, "shape1"), c(1, Inf))
  ## computations
  betaAW <- shape1 * (1 - QAW) / QAW
  out <- log(1 - obs[, "Y"])
  for (int in 1:shape1) {
    out <- out + 1/(int - 1 + betaAW)
  }
  out <- - out * shape1 * (1 - QAW) / QAW * 10 * sqrt(obs[, "W"]) * obs[, "A"]
  ## no need to center given how we will use it
  return(out)
}

vars <- eic(obs_from_another_experiment, psi = 59/300) *
  sigma0_draw_from_another_experiment(obs_from_another_experiment)
sd_hat <- sd(vars)
(slope_hat <- mean(vars))

## [1] 1.362463

(slope_CI <- slope_hat + c(-1, 1) * qnorm(1 - alpha / 2) * sd_hat / sqrt(B))

## [1] 1.357179 1.367746
```

Equal to 1.349 (rounded to three decimal places), the first numerical approximation `slope_approx` is not too off.

2.12 Double-robustness The efficient influence curve $D^*(P)$ at $P \in \mathcal{M}$ enjoys another remarkable property: it is double-robust. Specifically, if we define for all $P' \in \mathcal{M}$

$$\text{Rem}_P(\bar{Q}', \bar{G}') \equiv \Psi(P') - \Psi(P) + E_P(D^*(P')(O)), \quad (9)$$

then the so called remainder term $\text{Rem}_P(\bar{Q}', \bar{G}')$ satisfies^{††}

$$\text{Rem}_P(\bar{Q}', \bar{G}')^2 \leq \|\bar{Q}' - \bar{Q}\|_P^2 \times \|(\bar{G}' - \bar{G})/\ell\bar{G}'\|_P^2. \quad (10)$$

In particular, if

$$E_P(D^*(P')(O)) = 0, \quad (11)$$

and *either* $\bar{Q}' = \bar{Q}$ *or* $\bar{G}' = \bar{G}$, then $\text{Rem}_P(\bar{Q}', \bar{G}') = 0$ hence $\Psi(P') = \Psi(P)$. In words, if P' solves the so called P -specific efficient influence curve equation (11) and if, in addition, P' has the same \bar{Q} -component or \bar{G} -component as P , then $\Psi(P') = \Psi(P)$ no matter how P' may differ from P otherwise. This property is useful to build consistent estimators of $\Psi(P)$.

However, there is much more to double-robustness than the above straightforward implication. Indeed, 9 is useful to build a consistent estimator of $\Psi(P)$ that, in addition, satisfies a central limit theorem and thus lends itself to the construction of confidence intervals.

Let $P_n^0 \in \mathcal{M}$ be an element of model \mathcal{M} of which the choice is data-driven, based on observing n independent draws from P . Equality 9 reveals that the statistical behavior of the corresponding *substitution* estimator $\psi_n^0 \equiv \Psi(P_n^0)$ is easier to analyze when the remainder term $\text{Rem}_P(\bar{Q}_n^0, \bar{G}_n^0)$ goes to zero at a fast (relative to n) enough rate. In light of 10, this happens if the features \bar{Q}_n^0 and \bar{G}_n^0 of P_n^0 converge to their counterparts under P at rates of which *the product* is fast enough.

2.13 Inference assuming \bar{G}_0 known, or not, first pass. Let O_1, \dots, O_n be a sample of independent observations drawn from P_0 . Let P_n be the corresponding empirical measure, *i.e.*, the law consisting in drawing one among O_1, \dots, O_n with equal probabilities n^{-1} .

Let us assume for a moment that we know \bar{G}_0 . This may be the case indeed if P_0 was a controlled experiment. Note that, on the contrary, assuming \bar{Q}_0 known would be difficult to justify.

```
Gbar <- attr(obs, "Gbar")

iter <- 1e3 # 1e3 eventually
```

Then, the alternative expression

$$\psi_0 = E_{P_0} \left(\frac{2A - 1}{\ell\bar{G}_0(A, W)} Y \right) \quad (12)$$

suggests to estimate ψ_0 with

^{††}For any (measurable) $f : \mathcal{O} \rightarrow \mathbb{R}$, we denote $\|f\|_P = E_P(f(O)^2)^{1/2}$.

$$\psi_n^b \equiv E_{P_n} \left(\frac{2A-1}{\ell \bar{G}_0(A, W)} Y \right) = \frac{1}{n} \sum_{i=1}^n \left(\frac{2A_i-1}{\ell \bar{G}_0(A_i, W_i)} Y_i \right). \quad (13)$$

Note how P_n is substituted for P_0 in (13) relative to (12).

It is easy to check that ψ_n^b estimates ψ_0 consistently, but this is too little to request from an estimator of ψ_0 . Better, ψ_n^b also satisfies a central limit theorem: $\sqrt{n}(\psi_n^b - \psi_0)$ converges in law to a centered Gaussian law with asymptotic variance

$$v^b \equiv \text{Var}_{P_0} \left(\frac{2A-1}{\ell \bar{G}_0(A, W)} Y \right),$$

where v^b can be consistently estimated by its empirical counterpart

$$v_n^b \equiv \text{Var}_{P_n} \left(\frac{2A-1}{\ell \bar{G}_0(A, W)} Y \right) = \frac{1}{n} \sum_{i=1}^n \left(\frac{2A_i-1}{\ell \bar{G}_0(A_i, W_i)} Y_i - \psi_n^b \right)^2. \quad (14)$$

Let us investigate how ψ_n^b behaves based on `obs`. Because we are interested in the *law* of ψ_n^b , the next chunk of code constitutes `iter = 1000` independent samples of independent observations drawn from P_0 , each consisting of n equal to `nrow(obs)/iter = 1000` data points, and computes the realization of ψ_n^b on all samples.

Before proceeding, let us introduce

$$\begin{aligned} \psi_n^a &\equiv E_{P_n} (Y|A=1) - E_{P_n} (Y|A=0) \\ &= \frac{1}{n_1} \sum_{i=1}^n \mathbf{1}\{A_i=1\} Y_i - \frac{1}{n_0} \sum_{i=1}^n \mathbf{1}\{A_i=0\} Y_i \\ &= \frac{1}{n_1} \sum_{i=1}^n A_i Y_i - \frac{1}{n_0} \sum_{i=1}^n (1 - A_i) Y_i, \end{aligned}$$

where $n_1 = \sum_{i=1}^n A_i = n - n_0$ is the number of observations O_i such that $A_i = 1$. It is an estimator of

$$E_{P_0}(Y|A=1) - E_{P_0}(Y|A=0).$$

We seize this opportunity to demonstrate numerically the obvious fact that ψ_n^a does not estimate ψ_0 .

```
psi_hat_ab <- obs %>% as_tibble() %>% mutate(id = 1:n() %>% iter) %>%
  mutate(lGAW = A * Gbar(W) + (1 - A) * (1 - Gbar(W))) %>% group_by(id) %>%
  summarize(est_a = mean(Y[A==1]) - mean(Y[A==0]),
            est_b = mean(Y * (2 * A - 1) / lGAW),
            std_b = sd(Y * (2 * A - 1) / lGAW) / sqrt(n()),
            clt_b = (est_b - psi_approx) / std_b) %>%
  mutate(std_a = sd(est_a),
         clt_a = (est_a - psi_approx) / std_a) %>%
```




Figure 2: Kernel density estimators of the law of two estimators of ψ_0 (recentered with respect to ψ_0 , and renormalized), one of them misconceived (a), the other assuming that \tilde{G}_0 is known (b). Built based on *iter* independent realizations of each estimator.

```
gather("key", "value", -id) %>%
extract(key, c("what", "type"), "([^-_]+)_([ab])") %>%
spread(what, value)

(bias_ab <- psi_hat_ab %>% group_by(type) %>% summarise(bias = mean(cldt)))

## # A tibble: 2 x 2
##   type    bias
##   <chr> <dbl>
## 1 a     1.46
## 2 b     0.0696

fig <- ggplot() +
  geom_line(aes(x = x, y = y),
    data = tibble(x = seq(-3, 3, length.out = 1e3),
      y = dnorm(x)),
    linetype = 1, alpha = 0.5) +
  geom_density(aes(cldt, fill = type, colour = type),
    psi_hat_ab, alpha = 0.1) +
  geom_vline(aes(xintercept = bias, colour = type),
    bias_ab, size = 1.5, alpha = 0.5)

fig +
  labs(x = expression(paste(sqrt(n/v[n]^{list(a, b)})*(psi[n]^{list(a, b)} - psi[0]))))
```

Let v_n^a be n times the empirical variance of the *iter* realizations of ψ_n^a . By the above chunk of code, the averages of $\sqrt{n/v_n^a}(\psi_n^a - \psi_0)$ and $\sqrt{n/v_n^b}(\psi_n^b - \psi_0)$ computed across the realizations of the two estimators are respectively equal to 1.461 and 0.07 (both rounded to three decimal places — see `bias_ab`). Interpreted as amounts of bias, those two quantities are represented by vertical lines in Figure 2. The red and blue bell-shaped curves represent the empirical laws of ψ_n^a and ψ_n^b (recentered with respect to ψ_0 , and renormalized) as estimated by kernel density estimation. The latter is close to the black curve, which represents the standard normal density.

2.14 Inference assuming \bar{G}_0 known, or not, second pass. At the beginning of Section 2.13, we assumed that \bar{G}_0 was known. Let us suppose now that it is not. The definition of ψ_n^b can be adapted to overcome this difficulty, by substituting an estimator of $\ell_{\bar{G}_0}$ for $\ell_{\bar{G}_0}$ in (13).

For simplicity, we consider the case that \bar{G}_0 is estimated by minimizing a loss function on a single working model, both fine-tune-parameter-free. By adopting this stance, we exclude estimating procedures that involve penalization (*e.g.* the LASSO) or aggregation of competing estimators (*via* stacking/super learning) – see Section 2.15. Defined in the next chunk of code, the generic function `estimate_G` fits a user-specified working model by minimizing the empirical risk associated to the user-specified loss function and provided data.

Comment on new structure of `estimate_G` and say a few words about `compute_lGhatAW`.

```
estimate_G <- function(dat, algorithm, ...) {
  if (!attr(algorithm, "ML")) {
    dat <- as.data.frame(dat)
    fit <- algorithm[[1]](formula = algorithm[[2]], data = dat)
    Ghat <- function(newdata) {
      newdata <- as.data.frame(newdata)
      predict(fit, newdata, type = "response")
    }
  } else {
    fit <- algorithm(dat, ...)
    Qhat <- function(newdata) {
      caret::predict.train(fit, newdata)
    }
  }
  return(Ghat)
}

compute_lGhatAW <- function(A, W, Ghat, threshold = 0.05) {
  dat <- data.frame(A = A, W = W)
  Ghat_W <- Ghat(dat)
  lGAW <- A * Ghat_W + (1 - A) * (1 - Ghat_W)
  pred <- pmin(1 - threshold, pmax(lGAW, threshold))
  return(pred)
}
```

Note how the prediction of any $\ell_{\bar{G}_0}(A, W)$ is manually bounded away from 0 and 1 at the last but one line of `compute_lGhatAW`. This is desirable because the *inverse* of each $\ell_{\bar{G}_0}(A_i, W_i)$ appears in the definition of ψ_n^b (13).

For sake of illustration, we choose argument `working_model_G_one` of function `estimate_G` as follows:

```
working_model_G_one <- list(
  model = function(...) {glm(family = binomial(), ...)},
  formula = as.formula(
    paste("A ~",
      paste(c("I(W^", "I(abs(W - 5/12)^"),
        rep(seq(1/2, 3/2, by = 1/2), each = 2),
        sep = "", collapse = ") + ")),
    ")")
  ))
attr(working_model_G_one, "ML") <- FALSE
working_model_G_one$formula
```

```
## A ~ I(W^0.5) + I(abs(W - 5/12)^0.5) + I(W^1) + I(abs(W - 5/12)^1) +
##      I(W^1.5) + I(abs(W - 5/12)^1.5)
```

In words, we choose the so called logistic (or negative binomial) loss function L_a given by

$$-L_a(f)(A, W) \equiv A \log f(W) + (1 - A) \log(1 - f(W)) \quad (15)$$

for any function $f : [0, 1] \rightarrow [0, 1]$ paired with the working model $\mathcal{F} \equiv \{f_\theta : \theta \in \mathbb{R}^5\}$ where, for any $\theta \in \mathbb{R}^5$, $\text{logit } f_\theta(W) \equiv \theta_0 + \sum_{j=1}^4 \theta_j W^{j/2}$. The working model is well specified: it happens that \bar{G}_0 is the unique minimizer of the risk entailed by L_a over \mathcal{F} :

$$\bar{G}_0 = \arg \min_{f_\theta \in \mathcal{F}} E_{P_0} (L_a(f_\theta)(A, W)).$$

Therefore, the estimator \bar{G}_n output by `estimate_G` and obtained by minimizing the empirical risk

$$E_{P_n} (L_a(f_\theta)(A, W)) = \frac{1}{n} \sum_{i=1}^n L_a(f_\theta)(A_i, W_i)$$

over \mathcal{F} consistently estimates \bar{G}_0 .

In light of (13), introduce

$$\psi_n^c \equiv \frac{1}{n} \sum_{i=1}^n \left(\frac{2A_i - 1}{\ell \bar{G}_n(A_i, W_i)} Y_i \right). \quad (16)$$

Because \bar{G}_n minimizes the empirical risk over a finite-dimensional and well-specified working model, $\sqrt{n}(\psi_n^c - \psi_0)$ converges in law to a centered Gaussian law. Let us compute ψ_n^c on the same `iter` = 1000 independent samples of independent observations drawn from P_0 as in Section 2.13:

```
learned_features_fixed_sample_size <-
  obs %>% as_tibble() %>%
  mutate(id = 1:n() %% iter) %>%
  nest(-id, .key = "obs") %>%
  mutate(Ghat = map(obs, ~ estimate_G(., algorithm = working_model_G_one))) %>%
  mutate(lGAW = map2(Ghat, obs, ~ compute_lGhatAW(.y$A, .y$W, .x)))

psi_hat_abc <-
  learned_features_fixed_sample_size %>%
  unnest(obs, lGAW) %>%
  group_by(id) %>%
  summarize(est = mean(Y * (2 * A - 1) / lGAW)) %>%
  mutate(std = sd(est),
         clt = (est - psi_approx) / std,
         type = "c") %>%
  full_join(psi_hat_ab)
```




Figure 3: Kernel density estimators of the law of three estimators of ψ_0 (recentred with respect to ψ_0 , and renormalized), one of them misconceived (a), one assuming that \bar{G}_0 is known (b) and one that hinges on the estimation of \bar{G}_0 (c). The present figure includes Figure 2 (but the colors differ). Built based on `iter` independent realizations of each estimator.

```
(bias_abc <- psi_hat_abc %>% group_by(type) %>% summarise(bias = mean(clt)))
```

```
## # A tibble: 3 x 2
##   type    bias
##   <chr>  <dbl>
## 1 a      1.46
## 2 b      0.0696
## 3 c      0.0435
```

Note how we exploit the independent realizations of ψ_n^c to estimate the asymptotic variance of the estimator with v_n^c/n . By the above chunk of code, the average of $\sqrt{n/v_n^c}(\psi_n^c - \psi_0)$ computed across the realizations is equal to 0.043 (rounded to three decimal places — see `bias_abc`). We represent the empirical laws of the recentred (with respect to ψ_0) and renormalized ψ_n^a , ψ_n^b and ψ_n^c in Figures 3 (kernel density estimators) and 4 (quantile-quantile plots).

```
fig +
  geom_density(aes(clt, fill = type, colour = type), psi_hat_abc, alpha = 0.1) +
  geom_vline(aes(xintercept = bias, colour = type),
             bias_abc, size = 1.5, alpha = 0.5) +
  xlim(-3, 4) +
  labs(x = expression(paste(sqrt(n/v[n]^{list(a, b, c)}) *
                           (psi[n]^{list(a, b, c)} - psi[0]))))

ggplot(psi_hat_abc, aes(sample = clt, fill = type, colour = type)) +
  geom_abline(intercept = 0, slope = 1, alpha = 0.5) +
  geom_qq(alpha = 1)
```

Figures 3 and 4 reveal that ψ_n^c behaves as well as ψ_n^b — but remember that we did not discuss how to estimate its asymptotic variance.

2.15 Exercises. The problems come within the context of Sections 2.13 and 2.14.




Figure 4: Quantile-quantile plot of the standard normal law against the empirical laws of three estimators of ψ_0 , one of them misconceived (a), one assuming that \tilde{G}_0 is known (b) and one that hinges on the estimation of \tilde{G}_0 (c). Built based on `iter` independent realizations of each estimator.

1. Compute a numerical approximation of $E_{P_0}(Y|A = 1) - E_{P_0}(Y|A = 0)$. How accurate is it?
2. Building upon the piece of code devoted to the repeated computation of ψ_n^b and its companion quantities, construct confidence intervals for ψ_0 of (asymptotic) level 95%, and check if the empirical coverage is satisfactory. Note that if the coverage was exactly 95%, then the number of confidence intervals that would contain ψ_0 would follow a binomial law with parameters `iter` and 0.95, and recall that function `binom.test` performs an exact test of a simple null hypothesis about the probability of success in a Bernoulli experiment against its three one-sided and two-sided alternatives.
3. The call to `compute_lGhatAW` makes predictions on the same data points as those exploited to learn \tilde{G}_0 by fitting the user-supplied working model. Why could that be problematic? Can you think of a simple workaround, implement and test it?
4. Discuss what happens when the dimension of the (still well-specified) working model grows. You could use the following chunk of code

```
powers <- ## make sure '1/2' and '1' belong to 'powers', eg
seq(1/4, 3, by = 1/4)
working_model_G_two <- list(
  model = function(...) {glm(family = binomial(), ...)},
  formula = as.formula(
    paste("A ~",
          paste(c("I(W^", "I(abs(W - 5/12)^"),
                rep(powers, each = 2),
                sep = "", collapse = ") + ")),
          ")")
  ))
attr(working_model_G_two, "ML") <- FALSE
```

play around with argument `powers` (making sure that 1/2 and 1 belong to it), and plot graphics similar to those presented in Figures 3 and 4.


5. Discuss what happens when the working model is mis-specified. You could use the following chunk of code:

```

transform <- c("cos", "sin", "sqrt", "log", "exp")
working_model_G_three <- list(
  model = function(...) {glm(family = binomial(), ...)},
  formula = as.formula(
    paste("A ~",
          paste("I(", transform, sep = "", collapse = "(W)) + "),
          "(W))")
  )
attr(working_model_G_three, "ML") <- FALSE
(working_model_G_three$formula)

```

```
## A ~ I(cos(W)) + I(sin(W)) + I(sqrt(W)) + I(log(W)) + I(exp(W))
```

6.  Drawing inspiration from (14), one may consider estimating the asymptotic variance of ψ_n^c with the counterpart of v_n^b obtained by substituting $\ell\bar{G}_n$ for $\ell\bar{G}_0$ in (14). By adapting the piece of code devoted to the repeated computation of ψ_n^b and its companion quantities, discuss if that would be legitimate.

2.16 Inference based on the estimation of \bar{Q}_0 . Comment on structure of `estimate_Q`, similar to that of `estimate_G`.

Demonstrating the inference of ψ_0 based on the estimation of \bar{Q}_0 (and of the marginal law of W). Once based on a (mis-specified) working model, and once based on a non-parametric algorithm.

```

estimate_Q <- function(dat, algorithm, ...) {
  if (!attr(algorithm, "ML")) {
    dat <- as.data.frame(dat)
    fit <- algorithm[[1]](formula = algorithm[[2]], data = dat)
    Qhat <- function(newdata) {
      newdata <- as.data.frame(newdata)
      predict(fit, newdata, type = "response")
    }
  } else {
    fit <- algorithm(dat, ...)
    Qhat <- function(newdata) {
      caret::predict.train(fit, newdata)
    }
  }
  return(Qhat)
}

compute_QhatAW <- function(Y, A, W, Qhat, blip = FALSE) {
  if (!blip) {
    dat <- data.frame(Y = Y, A = A, W = W)
    pred <- Qhat(dat)
  } else {
    pred <- Qhat(data.frame(A = 1, W = W)) - Qhat(data.frame(A = 0, W = W))
  }
  return(pred)
}

working_model_Q_one <- list(
  model = function(...) {glm(family = binomial(), ...)},

```

```

formula = as.formula(
  paste("Y ~ A * (",
    paste("I(W~", seq(1/2, 3/2, by = 1/2), sep = "", collapse = ") + "),
    ")))")
))
attr(working_model_Q_one, "ML") <- FALSE
working_model_Q_one$formula

## Y ~ A * (I(W^0.5) + I(W^1) + I(W^1.5))

## k-NN
kknns_algo <- function(dat, ...) {
  args <- list(...)
  if ("Subsample" %in% names(args)) {
    keep <- sample.int(nrow(dat), args$Subsample)
    dat <- dat[keep, ]
  }
  caret::train(Y ~ I(10*A) + W, ## a tweak
    data = dat,
    method = "kknns",
    verbose = FALSE,
    ...)
}
attr(kknns_algo, "ML") <- TRUE
kknns_grid <- expand.grid(kmax = 5, distance = 2, kernel = "gaussian")
control <- trainControl(method = "cv", number = 2,
  predictionBounds = c(0, 1),
  allowParallel = TRUE)

learned_features_fixed_sample_size <-
  learned_features_fixed_sample_size %>% # head(n = 100) %>%
  mutate(Qhat_d = map(obs, ~ estimate_Q(., algorithm = working_model_Q_one)),
    Qhat_e = map(obs, ~ estimate_Q(., algorithm = kknns_algo,
      trControl = control,
      tuneGrid = kknns_grid))) %>%

  mutate(blip_QW_d = map2(Qhat_d, obs,
    ~ compute_QhatAW(.y$Y, .y$A, .y$W, .x, blip = TRUE)),
    blip_QW_e = map2(Qhat_e, obs,
      ~ compute_QhatAW(.y$Y, .y$A, .y$W, .x, blip = TRUE)))

psi_hat_de <-
  learned_features_fixed_sample_size %>%
  unnest(blip_QW_d, blip_QW_e) %>%
  group_by(id) %>%
  summarize(est_d = mean(blip_QW_d),
    est_e = mean(blip_QW_e)) %>%
  mutate(std_d = sd(est_d),
    std_e = sd(est_e),
    clt_d = (est_d - psi_approx) / std_d,
    clt_e = (est_e - psi_approx) / std_e) %>%
  gather("key", "value", -id) %>%
  extract(key, c("what", "type"), "([~_]+)_([de])") %>%
  spread(what, value)


```

```
(bias_de <- psi_hat_de %>% group_by(type) %>% summarize(bias = mean(clt)))
```

```
## # A tibble: 2 x 2
##   type    bias
##   <chr> <dbl>
## 1 d      0.304
## 2 e      0.0881
```

```
fig <- ggplot() +
  geom_line(aes(x = x, y = y),
    data = tibble(x = seq(-3, 3, length.out = 1e3),
      y = dnorm(x)),
    linetype = 1, alpha = 0.5) +
  geom_density(aes(clt, fill = type, colour = type),
    psi_hat_de, alpha = 0.1) +
  geom_vline(aes(xintercept = bias, colour = type),
    bias_de, size = 1.5, alpha = 0.5)

fig +
  labs(x = expression(paste(sqrt(n/v[n]^{list(d, e)})*(psi[n]^{list(d, e)} - psi[0]))))
```



No that bad! Yet, we know that \sqrt{n} times bias is bound to increase with sample size. To see this, check out the next chunks of code.

```
sample_size <- c(1e3, 4e3, 9e3)
block_size <- sum(sample_size)

label <- function(xx, sample_size = c(1e3, 2e3)) {
  by <- sum(sample_size)
  xx <- xx[seq_len((length(xx) %/% by) * by)] - 1
  prefix <- xx %/% by
  suffix <- findInterval(xx %/% by, cumsum(sample_size))
  paste(prefix + 1, suffix + 1, sep = "_")
}

learned_features_varying_sample_size <- obs %>% as.tibble %>%
  head(n = (nrow(.) %/% block_size) * block_size) %>%
  mutate(block = label(1:nrow(.), sample_size)) %>%
  nest(-block, .key = "obs")
```

First, we cut the data set into independent sub-data sets of sample size n in $\{1000, 4000, 9000\}$. Second, we infer ψ_0 as shown two chunks earlier. We thus obtain 71 independent realizations of each estimator derived on data sets of 3, increasing sample sizes.

```
learned_features_varying_sample_size <-
  learned_features_varying_sample_size %>%
  mutate(Qhat_d = map(obs, ~ estimate_Q(., algorithm = working_model_Q_one)),
         Qhat_e = map(obs, ~ estimate_Q(., algorithm = kknn_algo,
                                       trControl = control,
                                       tuneGrid = kknn_grid))) %>%

  mutate(blip_QW_d = map2(Qhat_d, obs,
                          ~ compute_QhatAW(.y$Y, .y$A, .y$W, .x, blip = TRUE)),
         blip_QW_e = map2(Qhat_e, obs,
                          ~ compute_QhatAW(.y$Y, .y$A, .y$W, .x, blip = TRUE)))

root_n_bias <- learned_features_varying_sample_size %>%
  unnest(blip_QW_d, blip_QW_e) %>%
  group_by(block) %>%
  summarize(clt_d = sqrt(n())*(mean(blip_QW_d) - psi_approx),
           clt_e = sqrt(n())*(mean(blip_QW_e) - psi_approx)) %>%
  gather("key", "value", -block) %>%
  extract(key, c("what", "type"), "([~_]+)_([de])") %>%
  spread(what, value) %>%
  mutate(block = unlist(map(strsplit(block, "_"), ~.x[2])),
         sample_size = sample_size[as.integer(block)])
```

The tibble called `root_n_bias` reports root- n times bias for all combinations of estimator and sample size. The next chunk of code presents visually our findings. Note how we include the realizations of the estimators derived earlier and contained in `psi_hat_de` (thus breaking the independence between components of `root_n_bias`, a small price to pay here).

```
root_n_bias <- learned_features_fixed_sample_size %>%
  mutate(sample_size = 1e3) %>%
  unnest(blip_QW_d, blip_QW_e) %>%
  group_by(id) %>%
  summarize(clt_d = sqrt(n()) * (mean(blip_QW_d) - psi_approx),
           clt_e = sqrt(n()) * (mean(blip_QW_e) - psi_approx),
           sample_size = sample_size[1]) %>% # because *fixed* sample size
  gather("key", "clt", -id, -sample_size) %>%
  extract(key, c("what", "type"), "([~_]+)_([de])") %>%
  mutate(block = "0") %>% select(-id, -what) %>%
  full_join(root_n_bias)

root_n_bias %>%
  ggplot() +
  stat_summary(aes(x = sample_size, y = clt,
                  group = interaction(sample_size, type),
                  color = type),
              fun.data = mean_se, fun.args = list(mult = 2),
              position = position_dodge(width = 250), cex = 1) +
  stat_summary(aes(x = sample_size, y = clt,
                  color = type),
              fun.y = mean,
              position = position_dodge(width = 250),
              geom = "polygon", fill = NA) +
```





Figure 5: Evolution of root- n times bias versus sample size for two inference methodology of ψ_0 based on the estimation of \bar{Q}_0 . Big dots represent the average biases and vertical lines represent twice the standard error.

```
geom_point(aes(x = sample_size, y = clt,
               group = interaction(sample_size, type),
               color = type),
           alpha = 0.1,
           position = position_dodge(width = 250)) +
labs(x = "sample size n",
     y = expression(paste(sqrt(n)*(psi[n]^{list(d, e)} - psi[0]))))
```

2.17 One-step estimation.

2.18 Targeted inference. For later...

```
working_model_Q_two <- list(
  model = function(...) {glm(family = binomial(), ...)},
  formula = as.formula(
    paste("Y ~ A * (",
          paste("I(W~", seq(1/2, 3, by = 1/2), sep = "", collapse = ") + ")",
                ")))")
)
attr(working_model_Q_two, "ML") <- FALSE

## xgboost based on trees
xgb_tree_algo <- function(dat, ...) {
  caret::train(Y ~ I(10*A) + W,
               data = dat,
               method = "xgbTree",
               trControl = control,
               tuneGrid = grid,
               verbose = FALSE)
}
attr(xgb_tree_algo, "ML") <- TRUE
xgb_tree_grid <- expand.grid(nrounds = 350,
                           max_depth = c(4, 6),
                           eta = c(0.05, 0.1),
```

```

        gamma = 0.01,
        colsample_bytree = 0.75,
        subsample = 0.5,
        min_child_weight = 0)

## nonparametric kernel smoothing regression
npreg <- list(
  label = "Kernel regression",
  type = "Regression",
  library = "np",
  parameters = data.frame(parameter =
    c("subsample", "regtype",
      "ckertype", "ckerorder"),
    class = c("integer", "character",
      "character", "integer"),
    label = c("#subsample", "regtype",
      "ckertype", "ckerorder")),
  grid = function(x, y, len = NULL, search = "grid") {
    if (!identical(search, "grid")) {
      stop("No random search implemented.\n")
    } else {
      out <- expand.grid(subsample = c(50, 100),
        regtype = c("lc", "ll"),
        ckertype =
          c("gaussian",
            "epanechnikov",
            "uniform"),
        ckerorder = seq(2, 8, 2))
    }
    return(out)
  },
  fit = function(x, y, wts, param, lev, last, classProbs, ...) {
    ny <- length(y)
    if (ny > param$subsample) {
      ## otherwise far too slow for what we intend to do here...
      keep <- sample.int(ny, param$subsample)
      x <- x[keep, ]
      y <- y[keep]
    }
    bw <- np::npregbw(xdat = as.data.frame(x), ydat = y,
      regtype = param$regtype,
      ckertype = param$ckertype,
      ckerorder = param$ckerorder,
      remin = FALSE, ftol = 0.01, tol = 0.01,
      ...)
    np::npreg(bw)
  },
  predict = function(modelFit, newdata, preProc = NULL, submodels = NULL) {
    if (!is.data.frame(newdata)) {
      newdata <- as.data.frame(newdata)
    }
    np::predict.npregression(modelFit, se.fit = FALSE, newdata)
  },

```

```

sort = function(x) {
  x[order(x$regtype, x$ckerorder), ]
},
loop = NULL, prob = NULL, levels = NULL
)

npreg_algo <- function(dat, ...) {
  caret::train(working_model_Q_one$formula,
    data = dat,
    method = npreg, # no quotes!
    verbose = FALSE,
    ...)
}
attr(npreg_algo, "ML") <- TRUE
npreg_grid <- data.frame(subsample = 100,
  regtype = "lc",
  ckertype = "gaussian",
  ckerorder = 4,
  stringsAsFactors = FALSE)

```